
PERPETUAL MOTION SQUAD
POCKETEXPENSE+

Final Project Report

Production-Grade Offline-First Financial Tracking System

Purvansh Sahu Yash Arya Aman Anand

TEAM	Perpetual Motion Squad
DEGREE	BSc (Hons) in Computer Science
INSTITUTION	BITS Pilani
PROJECT GUIDE	Mr. Murali Raghu
DATE	August 2026
DOCUMENT	Final Project Report

Contents

1. Executive Summary

2. Introduction

2.1 Problem Statement · 2.2 Objectives · 2.3 Scope

3. Literature Survey

4. System Architecture

4.1 Overview · 4.2 Frontend · 4.3 Backend · 4.4 Sync Engine

5. Feature Documentation (12 features)

6. Database Design

7. API Reference

8. Implementation Details

8.1 SMS Confidence Scoring · 8.2 Budget Alerts · 8.3 Sync · 8.4 Recurring · 8.5 Anomalies

9. Testing Strategy

10. Security Implementation

11. Challenges and Solutions

12. Future Enhancements

13. Conclusion

14. References

15. Appendices

1. Executive Summary

PocketExpense+ is a full-stack financial tracking application built on an offline-first architecture. It delivers twelve features spanning JWT authentication, expense management with offline synchronisation, category budgeting with threshold alerts, an idempotent recurring-transaction engine, aggregation-driven analytics with anomaly detection, receipt scanning with on-device OCR, and Android SMS-based automatic expense detection with measured **99% accuracy** and a **0% false-positive rate**.

The system comprises a React Native (Expo SDK 54) frontend and a Node.js/Express REST API backed by MongoDB. All client state is persisted locally via a Redux middleware persistence layer, so the application remains fully usable offline. A resilient synchronisation engine propagates local changes using tombstones, exponential backoff, and paged pulls.

Quality was a first-class deliverable: **198 automated tests** (96 backend, 102 frontend) with ~82% statement coverage on the backend, an SMS parser validated against 200 curated messages, and load-tested API targets of $p95 < 500$ ms and $p99 < 1000$ ms.

Headline results: 198/198 tests passing · 99% SMS accuracy · 0% false positives · ~82% backend coverage · 12/12 features complete · approved for production deployment.

2. Introduction

2.1 Problem Statement

Expense-tracking applications share two failure modes: they demand constant manual data entry, and they stop working without connectivity. Connected aggregators (Mint, Plaid-based apps) automate entry but require surrendering banking credentials, demand connectivity, and carry trust and regulatory risk — Mint's 2023 shutdown destroyed years of user financial history. PocketExpense+ removes the manual-entry burden on Android by parsing transaction SMS *locally on-device*, and remains fully functional offline.

2.2 Objectives

1. Deliver a cross-platform expense tracker that works offline-first.
2. Automate transaction capture on Android via local SMS parsing (target: $\geq 95\%$ accuracy, $\leq 3\%$ false positives).
3. Provide category budgeting with proactive threshold alerts.
4. Automate recurring transactions through an idempotent background processor.
5. Provide analytical insight via MongoDB aggregation pipelines.
6. Maintain a production-grade backend: JWT, validation, rate limiting, logging, graceful shutdown.
7. Reach a demonstrable quality bar: 198 automated tests and load-tested performance.

2.3 Scope

In scope: authentication; expense CRUD with offline sync and bulk delete; category budgets with alerts; recurring processing; filtering and pagination; insights; receipt OCR prefill; Android SMS detection with auto-add and undo; CSV/PDF export; local and in-app notifications; theming; production hardening.

Out of scope (§12): multi-currency, bank-credential aggregation, web application, cloud push, cloud OCR.

2.4 Project Team

This project was developed by the **Perpetual Motion Squad** — Purvansh Sahu, Yash Arya, and Aman Anand — as the final project for the BSc (Hons) in Computer Science at **BITS Pilani**, under the guidance of **Mr. Murali Raghu**. The work spans mobile application development, backend engineering, data modelling, and quality assurance, with each member contributing to design, implementation, and validation.

2.5 Deliverables

#	Deliverable	Document
1	Final project report	This document
2	Test case documentation	02-Test-Cases.pdf
3	Validation report	03-Validation-Report.pdf

4	Plagiarism compliance	04-Plagiarism-Compliance.pdf
5	User manual	05-User-Manual.pdf
6	Installation guide	06-Installation-Guide.pdf
7	Final presentation	07-Final-Presentation.pdf

3. Literature Survey

3.1 Personal Finance Applications

The personal-finance landscape is dominated by two architectures. Manual ledger applications such as basic trackers offer simplicity and offline reliability but place the entire data-entry burden on the user; churn studies consistently identify manual entry as the primary abandonment driver. Connected aggregators such as Mint and Plaid-based applications automate entry by connecting to banking APIs, but introduce three structural problems: (1) trust — users must surrender banking credentials to a third party; (2) fragility — Mint's 2023 shutdown destroyed years of user financial history; and (3) regulatory complexity — aggregating financial data triggers banking-regulator requirements that small teams cannot easily satisfy.

PocketExpense+ deliberately avoids the aggregator model. It automates entry from SMS transaction notifications — data the user already possesses — and does so entirely on-device, requiring no banking credentials and no external data partner.

3.2 Offline-First Architectures

Offline-first design is grounded in the CAP theorem (Brewer, 2000): during a network partition a distributed system must choose between availability and consistency. The offline-first approach chooses availability — "local mutations always succeed, and consistency is restored on reconnection." The literature establishes three core mechanisms that PocketExpense+ implements:

- **Tombstones** — logical delete markers that prevent deleted records from resurrecting during synchronisation, formalised in distributed-systems literature (DeCandia et al., 2007, Dynamo).
- **Client-side identifiers** — locally generated UUIDs that remain stable across synchronisation, allowing the server-assigned `_id` to be attached later without breaking references.
- **Idempotent reconciliation** — the sync endpoint returns `localId → serverId` mappings so the client can reconcile state deterministically even under retries.

These mechanisms are the difference between an app that "survives offline" (displays cached data) and one that is genuinely offline-first (fully mutable while disconnected).

3.3 SMS Parsing Approaches

Approach	Strengths	Weaknesses
Exact template matching per bank	High precision on known banks	Brittle; requires a large maintained library of exact templates
Machine-learning / NLP classifier	Generalises to unseen formats	Needs a labelled corpus; opaque failures; bundled model artefact
Heuristic regex + weighted confidence	Deterministic, testable, zero-dependency, inspectable	Needs careful pattern coverage

The project adopts the third approach: a five-stage pipeline (normalise → sender detection against 40+ bank patterns → classify → extract with strict decimal parsing → weighted confidence scoring). This is the approach that made a measured 99% accuracy with 0% false positives achievable on a 200-sample corpus.

3.4 Budgeting and Behavioural Finance

Behavioural-finance research shows that proactive, interruptive budgeting signals outperform retrospective reporting. A notification delivered at the moment a threshold is approached carries decision value; a report delivered a month later carries none. This motivates the budget-alert architecture: threshold logic fires local notifications at 80% and 100% of budget usage, once per budget per threshold per month. The alert is configured by the user (50/75/80/90% warning options) and is deliberately anti-flooding: a persisted `fired` key set prevents repeated alerts for the same threshold within a month.

3.5 Anomaly Detection in Finance

Anomaly detection in personal finance spans statistical methods (z-score, moving averages) and machine-learning methods (isolation forests, autoencoders). For a mobile application with sparse per-user data, the z-score method is attractive: it is interpretable, deterministic, and testable with zero dependencies. PocketExpense+ flags expenses exceeding two standard deviations from the user's mean, computed via MongoDB aggregation — no in-memory processing, and no model to train or ship.

Personal finance applications. Two architectures dominate: manual ledgers (reliable offline, high manual burden) and connected aggregators (automated but credential-heavy, fragile, and regulatorily complex). PocketExpense+ avoids both extremes by automating entry from SMS notifications entirely on-device.

Offline-first architectures. Grounded in the CAP theorem, offline-first design chooses availability — "local mutations always succeed; consistency is restored on reconnection." Three mechanisms are implemented: **tombstones** (DeCandia et al., 2007) to prevent deleted records resurrecting, **client-generated UUIDs** (`localId`) stable across sync, and **idempotent reconciliation** returning `localId → serverId` mappings.

SMS parsing. Three approaches exist: exact per-bank template matching (brittle), ML/NLP classifiers (require labelled corpora, opaque, bundled model artefacts), and heuristic regex + weighted confidence (deterministic, testable, zero-dependency). The project adopts the third: a five-stage pipeline achieves 99% accuracy and 0% false positives on a 200-sample corpus.


Budgeting and behavioral finance. Interruptive, proactive budget signals carry decision value at the moment a threshold is approached; retrospective reports carry none. This motivates alert delivery at 80% and 100% thresholds.

Anomaly detection. Z-score methods are interpretable, deterministic, and testable with zero dependencies — preferable to ML approaches for sparse per-user financial data. Expenses exceeding two standard deviations are flagged via aggregation.

4. System Architecture

4.1 Overview

The system uses a two-tier architecture with an intentionally thick client that owns local persistence, sync, SMS parsing, OCR orchestration, budget alerts, and analytics visualisation. The server is a thin, secure REST API for persistence, aggregation, recurring processing, and authentication.



4.2 Frontend Architecture

React Native 0.81.5 on Expo SDK 54; navigation via Expo Router 6 (typed routes); state via Redux Toolkit with seven slices, each owning a distinct concern. Persistence lives in `persistMiddleware` so reducers stay pure and testable.

Slice	State	Key Actions
auth	user, token, isAuthenticated, isLoading	hydrateAuth, login, logout, updateBudgetLimit
expenses	items, pendingQueue, tombstones, retry, totals	add/update/deleteExpense, deleteExpenses (bulk), markAsSynced, markDeleteSynced, queueRetry, applyServerSnapshot
budgets	items, loading, error	fetchAll, create, update, delete
insights	data, loading, error	fetchAdvanced
sync	isOnline, isSyncing, pendingCount, lastSyncTime, syncError	setOnlineStatus, setSyncing, setPendingCount, setLastSyncTime, setSyncError
sms	isEnabled, permissionStatus, lastDetectedTransaction, autoAdd state	enable/disableSmsDetection, setDetectedTransaction, setAutoAddEnabled, recordAutoAdded
notifications	items (feed, capped at 50)	addNotification, markRead, markAllRead, clearNotifications

4.3 Backend Architecture

Strict layering: `Routes → Middleware → Controllers (thin) → Services (logic) → Models`. Controllers contain no business logic; services own validation, aggregation, and business rules; models define schemas, indexes, virtuals, and hooks. Every response uses a uniform envelope: `{ success, message, data }`.

4.4 Offline-First Sync Engine

```
User adds expense → Redux store (optimistic) → AsyncStorage
                    |
                    Is online? — no → pendingQueue (sync later)
                    | yes
                    POST /api/expenses/sync → server reconciliation
                    |
                    localId → serverId mappings applied
```

- **Tombstones** — offline deletes are tracked and propagated so deletions never resurrect.
- **Exponential backoff** — retries at `2s → 4s → ... → 5min` (capped), stored as `nextAttemptAt`.
- **Paged pull** — snapshots fetched at 100 records/page, hard cap of 20 pages.
- **Force sync** — pull-to-refresh bypasses the backoff window.
- **Delete retry cap** — queued deletes give up after 5 attempts (`MAX_DELETE_ATTEMPTS`).

5. Feature Documentation

5.1 Authentication

Purpose: secure identity with stateless JWT sessions.

Implementation: registration validates name, email (regex), and password (min 6 chars); passwords hashed with bcrypt (10 rounds) and excluded from queries (`select: false`); JWT expires per `JWT_EXPIRE` (7d); an Axios interceptor injects the token on every request and 401 responses trigger automatic logout; profile updates support name, budgetLimit, currency, and avatar.

User flow: register → login → token stored in AsyncStorage → interceptor attaches `Authorization: Bearer <token>` → protected routes served; invalid or expired tokens redirect to login.

5.2 Expense Management + Offline Sync + Bulk Delete

Purpose: core tracking with validation and offline guarantees.

Implementation: each expense carries a client `localId` and `syncStatus` (synced/pending/conflict). New expenses write optimistically to the store and AsyncStorage; when online, the pending queue bulk-syncs via `POST /api/expenses/sync` and the server returns `localId → serverId` mappings. A NetInfo listener triggers sync on reconnect. Bulk delete uses a selection mode; offline deletes are tracked via tombstones. Supports 20 categories and 6 payment methods.

User flow: tap + → amount → category grid → payment method → date → save; the entry appears immediately whether online or not, and the sync badge tracks pending changes until uploaded.

5.3 Category Budgets + Threshold Alerts

Purpose: proactive per-category control with real-time progress feedback.

Implementation: monthly budgets per category with a compound unique index; Mongoose virtuals compute `percentageUsed` and `remainingAmount`; `totalSpent` recalculated via aggregation on every expense change. Pure threshold logic (`findCrossings`) fires warnings at ≥80% (configurable 50/75/80/90) and exceeded at ≥100%, once per budget per threshold per month, with side effects isolated in `budgetAlerts.ts`.

User flow: Account → Category Budgets → + → category and limit; progress bars colour-coded; local notifications at the warning and exceeded thresholds.

5.4 Recurring Transactions

Purpose: automate repeating transactions (rent, subscriptions, salary).

Implementation: frequency daily/weekly/monthly with `nextRunDate` computed on creation; `node-cron` runs hourly and finds due records via the `{isRecurring, nextRunDate}` index; for each due record it clones the template, advances `nextRunDate`, and sets `lastProcessedDate`. The `lastProcessedDate` guard provides idempotency under double-fire; recurring entries trigger budget recalculation.

User flow: add a transaction → toggle "Make this recurring" → choose frequency; entries generate automatically thereafter and are labelled "(recurring)".

5.5 Filtering + Pagination

Implementation: `GET /api/expenses` supports page, limit (max 100), category, type, date range, sort. Count and data queries run in parallel; six indexes cover common patterns; the `FilterModal` exposes the full set.

5.6 Smart Insights

Purpose: aggregate transaction data into decision-useful insight.

Implementation: basic insights (`GET /api/expenses/insights`) give monthly totals and category breakdown; advanced insights (`GET /api/insights/advanced`) compute 6-month growth, top-3 categories, weekday/weekend, and z-score anomalies — all via MongoDB aggregation. Home-screen stats compute locally (`src/utils/stats.ts`) for offline consistency.

User flow: Analytics tab for charts; Account → Smart Insights for the advanced analysis.

5.7 Receipt Scanning + OCR

Purpose: reduce manual entry by prefilling the expense form from a photographed receipt.

Implementation: `expo-camera` capture → optional ML Kit OCR provider (dev build; Expo Go falls back to photo attach + manual amount) → heuristic parser extracting amount via weighted total labels (grand total > total payable > total > amount/paid), merchant, and date, returning confidence and candidate amounts with ₹1–₹10,00,000 guardrails. The parser prefers `null` over a confidently wrong number.

User flow: + → Scan a receipt → photograph → fields prefilled → verify → save.

5.8 Android SMS Detection + Auto-Add

Implementation: five-stage pipeline (normalise → sender detect → classify → extract → score) runs fully on-device with 40+ bank sender patterns and 100+ keyword→category mappings. Dedup via LRU cache (200 entries) and a 2-minute guard. Routing: ≥ 0.75 opens the modal; 0.50–0.74 low-confidence; < 0.50 ignored; unrecognised sender below 0.65 always ignored. Auto-add ≥ 0.9 (configurable 0.8/0.9/0.95) logs immediately with a 6-second Undo toast.

5.9 Export (CSV / PDF)

Purpose: let users take their data out in portable formats.

Implementation: period presets (This month / Last month / Last 3 months / This year / All). CSV follows RFC 4180 escaping with signed amounts (expenses negative, income positive) so the column sums to the balance; PDF renders an HTML template via `expo-print` with escaping and a summary block (total expense/income/balance, by-category table); both open the native share sheet. All formatting logic is pure and unit-tested (`exportFormat.ts`).

User flow: Account → Export Data → choose period → CSV or PDF → share sheet.

5.10 Notifications

Implementation: `expo-notifications` local notifications with an Android channel; an Expo Go guard degrades gracefully; prefs persisted (enabled, warnThreshold default 80, notifyOnExceed, notifyOnAutoAdd); in-app feed capped at 50.

5.11 Theme System

Purpose: full light/dark/system theming without component branching.

Implementation: two palettes with identical key sets so components never branch on mode; modes light/dark/system persisted in AsyncStorage with an `isReady` gate preventing a scheme flash; `makeStyles()` memoizes `StyleSheet.create` per palette enabling runtime switches without teardown; dark-aware shadows via `elevation(isDark)`.

User flow: Account → Appearance → Light / Dark / Match System; the change is instant.

5.12 Production Hardening


Implementation: Joi validation everywhere; centralized `AppError` handler with production-safe responses; Helmet, mongo-sanitize, hpp, rate limiting (100/15min API, 20/15min auth), CORS allowlist, 10MB JSON limit; versioned routes; Winston logging; graceful shutdown.

5.13 Design Decisions and Trade-offs

Decision	Alternative	Rationale
Joi over express-validator	express-validator (installed)	Declarative, composable schemas that pair well with a service layer
Hourly cron vs event-driven	Bull/BullMQ job queue	Simpler at current scale; clear upgrade path
MongoDB aggregation for insights	Pre-computed materialised views	Keeps data fresh; ~200 ms latency is acceptable; Redis can be added
AsyncStorage offline-first	SQLite / WatermelonDB	Simpler for this data volume; WatermelonDB for >10K records
Z-score anomaly detection	Isolation Forest / ML	Interpretable, deterministic, testable, zero dependencies
Regex-based SMS parsing	ML/NLP model	Deterministic, testable, 99% accuracy on 200 samples
Confidence scoring over binary	Hard pass/fail	0% false positives while catching 99% of real transactions
LRU dedup cache (200 entries)	Bloom filter / server-side dedup	Client-only, no network, simple eviction
Persistence in Redux middleware	Persist inside reducers	Keeps reducers pure and unit-testable
Auto-add above 0.9 confidence	Always confirm first	Catches the common case with zero taps; Undo makes it reversible
Optional ML Kit OCR (dev build)	Cloud OCR always	Free + fully offline; graceful fallback keeps Expo Go usable

Local HTML+print PDF export	Native PDF libraries	Zero-dependency, one reusable template, bundle-safe
-----------------------------	----------------------	---

6. Database Design



6.1 Schemas (key constraints)

Collection	Fields / Constraints
User	name (req), email (unique, regex), password (min 6, bcrypt, select:false), budgetLimit (0), currency (INR), avatar
Expense	user (FK), amount (req), type (expense income), category (20 enum), description, paymentMethod (6 enum), date, isRecurring, frequency (daily weekly monthly null), nextRunDate, lastProcessedDate, localId, syncStatus (synced pending conflict)
Budget	user (FK), category (20 enum), amount (≥ 0), month (1-12), year, totalSpent; virtuals $percentageUsed = (totalSpent / amount) \times 100$, $remainingAmount = \max(amount - totalSpent, 0)$

7. API Reference

Method	Endpoint	Auth	Description
POST	/api/auth/register	Public	Register
POST	/api/auth/login	Public	Login (JWT + user)
GET/PUT	/api/auth/profile	Private	Get / update profile
GET	/api/expenses	Private	List (paginated + filtered)
POST	/api/expenses	Private	Create expense
GET/PUT/DELETE	/api/expenses/:id	Private	Read / update / delete
POST	/api/expenses/sync	Private	Bulk sync offline data
GET	/api/expenses/insights	Private	Basic monthly insights
GET/POST	/api/budgets	Private	List / create budgets
PUT/DELETE	/api/budgets/:id	Private	Update / delete budget
GET	/api/insights/advanced	Private	Growth, top categories, weekday/weekend, anomalies
GET	/api/health	Public	Health check

7.1 Example — Register

```
POST /api/auth/register
{ "name": "Purvansh", "email": "purvansh@example.com", "password": "secure123" }

→ { "success": true, "message": "Registration successful",
    "data": { "user": { "_id": "65f...", "name": "Purvansh", "email": "purvansh@example.com",
                     "budgetLimit": 0, "currency": "INR" },
              "token": "eyJhbGciOiJIUzI1NiIsInR5cCI6IWR5bmlzInQ=" } }
```

7.2 Example — Create Expense

```
POST /api/expenses
Authorization: Bearer <token>
{ "amount": 250, "category": "food", "paymentMethod": "upi",
  "isRecurring": true, "frequency": "monthly" }

→ { "success": true, "message": "Expense created successfully",
    "data": { "_id": "65f...", "amount": 250, "type": "expense",
              "category": "food", "paymentMethod": "upi",
              "nextRunDate": "2026-04-03T00:00:00.000Z",
              "syncStatus": "synced", "date": "2026-03-03T00:00:00.000Z" } }
```

7.3 Example — List Expenses (Filtered)

```
GET /api/expenses?category=food&type=expense&startDate=2026-01-01&endDate=2026-03-31&page=1&limit=20&sort=amount

→ { "success": true, "message": "Expenses fetched",
    "data": { "expenses": [ { ... } ], "totalPages": 5, "currentPage": 1, "totalItems": 98 } }
```

7.4 Example — Advanced Insights

GET /api/insights/advanced

```
→ { "monthlyGrowthRate": { "rates": [ { "month": 2, "year": 2026,
    "total": 5200, "growthRate": 30.5 } ], "averageGrowthRate": 15.25 },
  "topCategories": [ { "_id": "food", "total": 3200, "count": 45, "avgAmount": 71.11 } ],
  "weekdayVsWeekend": { "weekday": { "total": 4500, "count": 60, "avgPerTransaction": 75 },
    "weekend": { "total": 1800, "count": 20, "avgPerTransaction": 90 },
    "comparison": "Weekend spending is 28.6% of total" },
  "anomalies": { "detected": true,
    "stats": { "mean": 120.5, "stdDev": 45.2, "totalAnalyzed": 150 },
    "items": [ { "amount": 5000, "category": "travel", "zScore": 4.82 } ] }
```

8. Implementation Details

8.1 SMS Confidence Scoring Engine

Signal	Weight	Penalty	Weight
Recognised bank sender	+0.25	Ambiguous keyword (fraud/suspicious)	-0.20
Valid amount parsed	+0.30	Multiple amounts in message	-0.15
Merchant extracted	+0.20	Suspicious formatting (URLs, punctuation)	-0.10
Transaction keyword	+0.15		
Account digits	+0.05		
Timestamp	+0.05		

Routing: $\geq 0.75 \rightarrow$ modal; $0.50-0.74 \rightarrow$ low-confidence banner; $< 0.50 \rightarrow$ ignore; unrecognised sender + $< 0.65 \rightarrow$ always ignore; auto-add at ≥ 0.9 with Undo. Guardrails: ₹1–₹10,00,000, strict decimal parsing, merchant ≥ 2 chars, min SMS length 15, sender gating below 0.65.

```
Incoming SMS (Android OS)
  | smsListener (singleton, 100ms debounce)
  ▼
1. normalizeMessage()    strip zero-width chars, normalise ₹/INR
2. detectSenderMetadata() match 40+ bank sender-ID patterns
3. classifyMessageType()  debit / credit / ignore
4. extractFields()        amount, merchant, account, timestamp
5. computeConfidence()    weighted 0.0–1.0 scoring
  |
  ▼
Deduplication: LRU cache (200) + 2-minute time-window guard
  |
  ▼
Category auto-detection: 100+ keyword → category mappings
  |
  ▼
Confidence routing → confirm modal / auto-add ( $\geq 0.9$ ) + Undo
```

8.2 Budget Alert Engine

`findCrossings` is pure and unit-tested in isolation; side effects (notifications, feed) live in `budgetAlerts.ts`. Warning fires at $\geq$ `warnThreshold` (default 80; options 50/75/80/90); exceeded at $\geq 100\%$. Each alert fires once per budget per threshold per month via a persisted `fired` key set reset on month rollover.

8.3 Sync Engine Algorithms

The sync engine implements several resilience algorithms:

- **Exponential backoff:** retry delays follow `2s → 4s → 8s → ...` capped at 5 minutes, stored as `nextAttemptAt`; pull-to-refresh forces a sync and bypasses the window.
- **Tombstones:** offline deletes append a tombstone record; on the next sync the server applies them and returns acknowledgement, preventing deleted records from resurrecting.

- **Paged pull:** snapshots are fetched at 100 records/page with a hard cap of 20 pages, so a flaky response cannot loop indefinitely.
- **Delete retry cap:** `MAX_DELETE_ATTEMPTS = 5`; after that the queued delete is dropped rather than retried forever.
- **Reconciliation:** the `/sync` endpoint returns `localId → serverId` mappings used to update the client store deterministically.

8.4 Recurring Transactions

Hourly cron finds `nextRunDate ≤ now` via the `{isRecurring, nextRunDate}` index; clones each template, advances the date, sets `lastProcessedDate`; the guard makes re-execution idempotent; budget recalculation is triggered.

8.5 Anomaly Detection

Z-score via aggregation: $z = (\text{value} - \text{mean}) / \text{stdDev}$; flags expenses beyond 2 standard deviations; returns amount, category, date, and z-score. Deterministic and testable.

9. Testing Strategy

All suites green: **backend 96/96** — **frontend 102/102**.

9.1 Backend (96 tests, 6 suites)

Layer	Tests	Coverage
Expense Service	15	91%
Budget Service	12	93%
Recurring Service	10	91%
Insight Service	11	91%
API Integration	17	100% routes
Edge Cases	31	zero data, concurrency, precision, timezone
Total	96	~82%

```
cd server
npm test           # all
npm run test:unit  # 48 service-level
npm run test:integration # 17 API-level
npm run test:edge  # 31 edge cases
npm run test:coverage
```

9.2 Frontend (102 tests, 6 suites)

Suite	Tests	Covers
tests/sms/smsParser.test.ts	42	pipeline stages, 200-sample accuracy ($\geq 95\%$, $FP \leq 3\%$), guardrails
tests/receipt/parseReceipt.test.ts	16	weighted total labels, merchant/date, confidence, negatives
tests/export/exportFormat.test.ts	16	CSV escaping/signing, windowing, HTML escaping
tests/budget/thresholds.test.ts	12	findCrossings, once-per-month dedup, custom thresholds
tests/expense/bulkDelete.test.ts	4	items + pendingQueue removal, no-op, tombstones
tests/stats/stats.test.ts	12	totalsForMonth, monthOverMonthChange
Total	102	

```
npm test           # all 102
npm run test:sms   # SMS parser suite only (42)
```

9.3 SMS Parser Accuracy

SMS PARSER ACCURACY REPORT

Total samples:	200	Accuracy:	99%
True positives:	158	False + rate:	0%
True negatives:	40	Avg confidence:	0.92
False positives:	0		
False negatives:	2		

Metric	Target	Actual
Detection accuracy	$\geq 95\%$	99%
False positive rate	$\leq 3\%$	0%
Debit detection rate	$\geq 95\%$	100%
Credit detection rate	$\geq 90\%$	95%
OTP rejection	100%	100%
Promo rejection	100%	100%
Avg confidence score	≥ 0.6	0.92

9.4 Performance Testing

```
cd server
npm install -g artillery
artillery run tests/performance/artillery.yml
```

- $p95 < 500 \text{ ms}$ · $p99 < 1000 \text{ ms}$ · 100 concurrent users · 500 requests/min
- SMS parser performance measured via `simulateBulkSms(count)` (dev-only).

10. Security Implementation

Control	Status
Helmet security headers (CSP, HSTS, X-Frame-Options, X-Content-Type-Options)	✓
NoSQL injection prevention (express-mongo-sanitize)	✓
HTTP parameter pollution prevention (hpp)	✓
Rate limiting: 100 req/15min API, 20 req/15min auth	✓
JWT verification on all private routes	✓
bcrypt password hashing (10 salt rounds); password <code>select:false</code>	✓
Joi validation on all endpoints; unknown fields stripped	✓
Stack traces hidden in production; CORS allowlist; 10MB body limit	✓
Graceful shutdown (SIGTERM/SIGINT)	✓
No SMS content logged or transmitted; parsing entirely local	✓
CSV formula-injection escaping; HTML escaping in PDF exports	✓

10.1 JWT Authentication Flow

1. User submits email + password.
2. Server verifies via bcrypt `matchPassword`.
3. Server signs a JWT with a 7-day expiry (`JWT_EXPIRE`).
4. Client stores the token in AsyncStorage and attaches it via the Axios interceptor.
5. 401 responses trigger automatic logout and credential clearing.

10.2 SMS Privacy Principles

1. **Local-only parsing** — no SMS content ever leaves the device.
2. **No raw SMS storage** — the body is processed in-memory and discarded; only structured fields are retained.
3. **No SMS logging** — debug logging stripped in production.
4. **User-controlled and revocable** — off by default; disabling clears the dedup cache and disables auto-add.
5. **Minimal permissions** — only `READ_SMS` and `RECEIVE_SMS`.
6. **Backend isolation** — no endpoint receives SMS content; entries arrive as ordinary expenses.
7. **Non-reversible fingerprints** — dedup hashes cannot reconstruct the original message.

11. Challenges and Solutions

Challenge	Solution
OCR / SMS need native modules absent from Expo Go	Provider isolation (<code>ocr.ts</code> , <code>smsListener.ts</code>) + graceful Expo Go fallback
expo-notifications drops remote support in Expo Go	Module import guarded; calls degrade to no-op
Duplicate recurring transactions if cron fires twice	Idempotency guard on <code>lastProcessedDate</code>
Offline deletes resurrecting after sync	Tombstones propagated on next sync
Flaky networks causing unbounded retry loops	Exponential backoff (5min cap), paged pulls (20 pages), delete retry cap
False positives writing wrong financial records	Weighted confidence + sender gating + confirmation modal + Undo
Theme switch tearing down component tree	Memoized <code>makeStyles()</code> per palette
Budget alert flooding	Once-per-budget-per-threshold-per-month dedup

12. Future Enhancements

1. Secure token storage (`expo-secure-store`)
2. Multi-currency with live exchange-rate conversion
3. Cloud push notifications (FCM / Expo push)
4. Bull/BullMQ job queue replacing `node-cron`
5. Redis cache for insights; read replicas; sharding by `user_id`
6. CDC for real-time budget recalculation
7. SMS ML model for >99.5% accuracy; cloud OCR endpoint
8. Web application (architecture supports it)

13. Conclusion

PocketExpense+ removes the two primary reasons expense applications fail — manual-entry fatigue and offline fragility. The Android SMS detection feature, with measured 99% accuracy and 0% false positives, demonstrates that dependable automation is achievable with deterministic, inspectable parsing. The offline-first architecture, with tombstones, backoff, and paged sync, provides genuine offline usability backed by a secure, rate-limited API.

The engineering quality bar was met and exceeded: 198 automated tests pass, backend coverage is ~82%, and load-test targets are defined. All twelve features are complete and documented. The system is ready for evaluation, demonstration, and production deployment subject to the deployment considerations in the Installation Guide.

13.1 Key Outcomes by Objective

Objective	Outcome
Offline-first operation	All core operations work with no connection; sync is automatic and resilient
SMS automation ≥95% accuracy	99% accuracy, 0% false positives on a 200-sample corpus
Proactive budgeting	80%/100% threshold alerts, once per budget per month
Recurring automation	Hourly idempotent processor, no duplicate entries
Analytical insight	Growth, top categories, weekday/weekend, z-score anomalies
Production-grade backend	JWT, validation, rate limiting, logging, graceful shutdown
Quality bar	198 tests passing, ~82% coverage, load-tested performance

14. References

- DeCandia, G. et al. (2007). Dynamo: Amazon's Highly Available Key-Value Store. SOSP.
- Brewer, E. A. (2000). Towards Robust Distributed Systems. PODC.
- React Native / Expo SDK / Redux Toolkit / Express / Mongoose / MongoDB official documentation.
- RFC 7519 (JWT) · RFC 4180 (CSV) · Helmet.js · Artillery load testing.

15. Appendices

Appendix A — Environment Variables

Variable	Default	Description
PORT	5001	Server port
HOST	0.0.0.0	Server host
NODE_ENV	development	Environment
MONGODB_URI	—	MongoDB connection string
JWT_SECRET	—	JWT signing secret (min 32 chars)
JWT_EXPIRE	7d	Token expiry
RATE_LIMIT_WINDOW_MS	900000	Rate limit window (15 min)
RATE_LIMIT_MAX	100	Max requests per window
CORS_ORIGIN	*	Allowed origins

Appendix B — Scalability Path

1. Redis cache for the insights endpoint (invalidate on expense write)
2. Bull/BullMQ job queue to replace node-cron for recurring transactions
3. MongoDB read replicas for analytics queries
4. Microservice extraction of the insight engine if query load grows
5. Change Data Capture for real-time budget recalculation
6. Sharding by `user_id` at single-node limits

Appendix C — Deployment

- **Backend (Render):** connect repo, set root `server`, env vars from `.env.example`, deploy via `render.yaml`. Docker: `docker build -t pocketexpense-api .` · `docker run -p 5001:5001 --env-file .env pocketexpense-api`.
- **Frontend (EAS):** profiles `development` (dev client), `preview` (internal), `production` (store-ready). SMS detection and OCR require a development or production build.
- **Note:** the application is distributed through Expo Go project links and EAS builds; it is not published on any public app store.

Appendix C — Design System

Typography: Poppins (400/500/600/700) across sizes xs (10) to hero (40). **Palette:** a violet-glassmorphism light theme and a dark theme, both defined in `src/theme/colors.ts` with identical key sets. Primary token `#8A64EB`, gradient `#C4A6FE → #8A64EB`, secondary lime `#C7F2A4`.

20 categories: Groceries, Travel, Car, Home, Insurance, Education, Marketing, Shopping, Internet, Water, Rent, Gym, Subscription, Vacation, Food, Entertainment, Salary, Freelance, Investment, Other.

6 payment methods: Cash, Credit Card, Debit Card, Bank Transfer, UPI, Other.

Appendix D — Project Structure

```
finalAssignment/
├─ app/                # Expo Router screens (auth, tabs, expense, settings)
├─ src/
│  ├─ components/      # common, expense, home, navigation, sms, sync, analytics
│  ├─ services/        # sms/, receipt/, syncEngine, budgetAlerts,
│  │                   # budgetThresholds, categoryDetector, exportFormat,
│  │                   # export, notifications, smsListener, smsPermission
│  ├─ store/           # apiClient, slices/, persistMiddleware
│  ├─ theme/           # colors, ThemeContext, makeStyles
│  └─ utils/           # formatters, id, stats
├─ tests/              # sms, receipt, export, budget, expense, stats
├─ server/
│  ├─ config/ controllers/ middleware/ models/ routes/
│  ├─ services/ validators/ tests/ docs/
│  ├─ app.js index.js jest.config.js Dockerfile render.yaml .env.example
└─ app.json eas.json jest.config.js tsconfig.json package.json
```

Guide Details

Project Guide	Prof. Murali Raghu
Institution	BITS Pilani
Degree	BSc (Hons) in Computer Science
Date	28 August 2026
Time	01:58 PM

Prof. Murali Raghu

28 August 2026

Signature

Date / Time